


Process View: Code Guardian Analysis Workflows



Legend:



Typical latency pattern:

- Cache hit: no new LLM call
- Cache miss: dominated by model latency
- RAG adds modest retrieval overhead

Key optimizations:

- Debouncing reduces repeated calls during typing
- Function-level scoping limits prompt size
- Cache helps on repeated snippets
- RAG retrieval depth is configurable

Privacy guarantees:

- All code analysis local (no exfiltration)
- LLM runs on localhost:11434
- Vector DB stored locally
- Optional CVE refresh: metadata only